

Papyr: Automatically Generating “Smart” Paper Tools for Low-resource Settings

ABSTRACT

Despite the recent push toward leveraging information and communication technologies (ICTs) for replacing paper-based tasks, there remain many barriers to designing deployable and appropriate technical solutions that replace paper. As a result, paper tools such as forms, charts, and graphs continue to be widely used in low-resource organizations working in developing regions. Unfortunately, existing paper tools are often designed to be out of reach for end-users and depend on intermediaries to input data and interpret information. We apply ideas from computability and human-computer interaction to design a system to automatically generate paper tools that provide immediate computation, visual feedback, and independence from both ICTs and intermediaries at the point of use. **Papyr** is a system that is easy to use and allows low-resource organizations to quickly transform data tracking or information dissemination requirements into printable “smart” paper tools. In this paper, we explore how paper tools are currently being used by organizations in Ghana, explore the design space for smart paper tools, and describe **Papyr**’s design and implementation.

1. INTRODUCTION

Even as computing devices proliferate and try to eliminate paper from a given ecosystem, it is undeniable that paper continues to be pervasive in professional and personal spaces. Paper is light, low-cost, familiar, accessible, easy to use, intuitive to manipulate, and convenient to distribute [20, 21, 12]. Moreover, in low-resource settings where technology infrastructure, such as stable electricity, cutting-edge computer hardware and software, and skilled human capital, is deficient, paper is an obvious choice because it can accomplish many of the same tasks that computers can, including tracking, recording, and look-up. Furthermore, generations of paper-based organizational and end-user tasks have ensured that national regulations seek compliance with paper-based record keeping rules. The enduring legacy of paper, especially amongst current paper-based workflows, provides

inertia against switching to more technically intensive replacements.

Still, paper is not without its constraints. When compared to information and communication technology (ICT) alternatives, paper is less malleable once data is inscribed, has less “computational output”, and is difficult to digitize. Electronic media that is designed and built to replace paper possesses an inverse profile of not only strengths but also constraints, unable to capitalize on paper’s more accessible and tangible properties [12]. Can paper itself be designed to overcome some of its limitations? Can paper, for instance, make its static content dynamic? Or, perhaps, even more ambitiously, can paper provide computational feedback to its end users?

There already exist many examples of “smart paper” that demonstrate how these lofty objectives are in fact possible. Paper graphs, for instance, can provide computational feedback. The nomograph has three scales to represent a three-variable equation. In one illustration, by plotting along these scales using just a pen, a user can determine their ideal body mass by locating their results within the “overweight”, “desirable”, or “underweight” categories, while accounting for any gender differences [25]. The partograph, another example, is a graphical tool that can provide predictive feedback to doctors or midwives throughout the progress of labour [14]. It is popular in low-resource settings, a cost-effective and accessible tool that when coupled with well-defined protocols and appropriate instruction has been shown to demonstrate distinct benefits [9, 26, 8, 17].

These examples of “smart” paper tools indicate a particularly compelling research problem of leveraging paper’s many useful properties and developing more standalone tools, validating their usefulness in low-resource settings where paper is cheaper, more familiar, and more accessible than its technological counterparts. Our study addresses this research problem and explores the “smart” paper space with the following contributions:

- Needs assessment of low-resource settings for micro-finance and health in Ghana through an appraisal of existing paper artifacts and understanding of information gaps. (Section 3).
- Development of “smart” paper tools that are pushing the boundaries of paper and exploring how paper can contribute to existing processes, especially with low-literate populations in mind (Section 3).
- Discovery of basic design principles that highlight the constraints of paper and user interaction with “smart” paper tools (Section 4).

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

DEV ’14 San Jose, California USA

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.




Figure 1: Spectrum between paper and technology.

- Construction of **Papyr**, a system for intermediaries from low-resource organizations to automatically generate “smart” paper tools by simply defining their tasks and their associated parameters. The intermediary can then select from suggestions based on resources available and capabilities of the end-user to print and construct the paper tools based on the resources available (Section 5).

2. RELATED WORK

There is a spectrum of innovations that attempt to bridge the gap between technology and paper and leverage the benefits of both, a spectrum based on the degree to which the innovation is integrated with technology, as described in Figure 1.

The tools that rest in the space between paper and technology most strongly integrated with technology are *paper-like technologies*. This includes a set of tablets and mobile devices, including PaperTab, PaperFold, and MorePhone, that are electronic devices that are as reconfigurable as paper itself [11, 10, 24]. The selling point of these technologies are their paper-like properties, including portability, maneuverability, and flexibility, a familiar look and feel to technology with the provision of the same capacities as a full-fledged tablet or phone. This reinforces the draw to include the benefits of paper in technological design, but does so in a fairly expensive way.

Other tools attempt to explore the role of paper documents as interfaces to the digital world, or the field of *paper computing* [13]. This includes systems that use paper to control technological environments, such as Digital Desk, the projection of graphics interfaces on a paper-based workspace, or PaperPoint, the control and annotation of PowerPoint presentations using printed slide handouts [15, 30, 22]. In addition, this includes tools that provide technological links to paper, such as printed visual markers (QR codes and barcodes); specialized devices for annotations and digitally inputted notes; computer vision based digitization of paper; and cloud computing [13, 2, 18]. However, these technologies often use paper as an aid to technology as opposed to utilizing paper’s beneficial properties on its own.

Paper can be used to model technology in paper prototyping or paper composites. *Paper prototyping* utilizes paper’s accessibility and ability to be easily manipulated to simulate and test the usability technological designs in a “wizard of oz” experimental structure [23, 3]. This shows a loose coupling of paper to technology in that technology is not used in the generation of the tool, but the tool is paradoxically dependent on technology because it has no power when used alone. On the other hand, *paper composites*, or combinations of paper with technological materials, such as the paper and conductive thread combination in pulp-based computing, can give these paper-based tools more power, but still only exist to simulate technology and are reliant on

technological components [5].

One of the problems with this type of relationship with technology and technological components in the workflow is cost and a reliance on tools and materials low-resource organizations do not necessarily have. However, there are some innovations that attempt to bridge the gap between existing paper materials with low-cost technological interfaces like mobile phones.

Looking specifically at some of the tools developed for low-resource settings, most are still tightly reliant on technology and would fall mostly in the *paper computing* category. This includes tools that recapture information that already exists on paper into the technological space, such as CAM, which uses a camera to record and store data from paper logs in rural microfinance groups in India; ODK Scan, which applies computer vision techniques on small “snippets” of the content for a single form field; and Shredder, a combination of computer vision, database, and crowdsourcing techniques to transmit and verify information from paper forms [16, 1, 7, 6, 4]. There are also tools that use technology to track the actual process for data input on paper-based forms such as Digital Slate, which uses the pen to paper writing motion to capture data in savings groups in India, and Partopen, which is a digital pen specifically developed for filling out the partograph in maternal labor [19, 27, 28]. And there are even tools that use a mix of technology and paper in the information gathering process itself, such as Local Ground, a barcoded geospatial surveying tool [29].

In addition to being tightly coupled with technology, these tools are used mainly for information collection, including both retroactive and active data capture. Most of these innovations do not give very much feedback to the user at the time of use, except perhaps PartoPen, and thus does not leverage one of the important abilities of computation.

We believe **Papyr** is an unusual but appropriate marriage between paper and technology for developing contexts: a computer is used to support the organization’s design of context-specific paper tools, but the produced paper tools require no ICT support or leverage at the end user. Moreover, **Papyr** provides feedback to the end-users, encouraging utilization and adherence to the processes.

3. EXPLORING PAPER

After reviewing related work and identifying real-world problems in the health and microfinance contexts in low-resource settings where paper is already deeply rooted and in constant use, we developed two sets of design iterations: one based solely on brainstorming sessions and the other based on in the field feedback. In Section 3.1, we describe designs that attempt to push the boundaries of paper to explore the range of computational and visual feedback that paper in and of itself could provide. Afterwards, in Section 3.2, we discuss designs that are based on a needs assessment to see how paper is being used in the health and microfinance contexts at field sites in Ghana and feedback from real end users.

3.1 Brainstorming Design Iterations

The primary goal of these preliminary brainstorming sessions was to push the boundaries of what “smart paper” could accomplish. In other words, how could we re-design paper in a way that makes it somehow better? Here, we describe a small subset of the “smart-paper” tools we designed.




Figure 2: Subtraction band described in Section 3.1.1.

3.1.1 The Subtraction Band

Use cases. One problem in the microfinance domain for low-literate users is keeping track of exactly how much is earned and where money is spent in a given day. We designed a tool that would help tackle one of the sources of difficulty for this problem, or doing addition and subtraction of expenditures more easily.

Proposed tool. The base is made of stiff paper, plastic, cardboard, or other material that could support two narrow slits through which a closed paper loop could be smoothly slid and rotated. There are two sections on the base, the top section is a basic scale of amounts with a predetermined interval and on the bottom is the band threaded through two slits. This band can be made with paper or waxed paper to ensure smoother sliding. A depiction can be seen in Figure 2.

An individual will place a sticker in each slot every time they earn money that day. To take a simple example, say an individual has earned cumulatively 10 cedi¹ today. The band would be marked up to the 10th slot. However, the individual now spends 3 cedis. They will remove the stickers up to the 3rd slot. Then, they can shift the band to the left such that the first filled slot moves back to the leftmost point, or the 1st slot. Depending on the resources available, the individual can opt to shade in the band instead of using stickers.

Lessons learned. This tool can be generalized from the basic use case to include any type of task that requires addition and subtraction.

3.1.2 The Addition By Matching Tool

Use cases. Low-income populations are typically engaged in the informal labour market, meaning their income inflows are irregular and variable, accumulated through a variety of financial instruments such as money earned through different jobs and money borrowed through different sources of credit. We decided to design a tool that can be used to compare income and expenses to determine when income-expenses can be met or matched versus when one exceeds the other as well as allow the user to flexibly update any changes in the projected values.

Proposed tool. Again, the base would be made of stiff paper, plastic, cardboard, or other material that would support two narrow slits through which a paper strip could be smoothly slid up or down. Figure 3 is a representation of this tool. The horizontal axes indicates where each separate strip would be strung and the vertical axes indicate the amounts accumulated on each strip.

If a user estimates that the incomes for the first day will

¹local currency corresponding to roughly 0.28 USD per cedi at the time of writing




Figure 3: Addition by matching described in Section 3.1.2.

be around 20 cedi, she or he would fill the cells in the income column for Day 1 up to the 20 cedi mark on the vertical axis. And then if the user estimates that on the second day, the income will be approximately 30 cedi, she or he would fill the cells up to the 30 cedi mark on the axis and then shift the band so that the stickers start after the filled income for the previous day, at the 20 cedi mark. In this case, the user can easily tell that over the two days, the income has been accumulated to 50 cedis. The user then can finish populating the rest of her or his projected income in a similar manner and continue in a similar manner for the expenses columns. In the example, the user can then see that over three days, she or he as earned and spent cumulatively 60 cedis. The user can also see that expenses exceed income on Day 1, income exceeds expenses on Day 2, and they match on Day 3.

While tracking incomes and expenses, if there is an unexpected 10 cedis of expense that is introduced on Day 1, then the successive stacks of expenses must be readjusted. This is achieved by sliding the paper strip for the expenses on the following day up so that the filled in units rest above the end of the filled units for the day that was altered, and continued through to all of the successive days.

Lessons learned. This tool can be generalized to calculate any cumulative sum and compare and match analogous variables.

3.1.3 Averages Tool

Use cases. Often times in villages, women are not actively tracking their menstrual cycle. Therefore, the woman does not necessarily know when ovulation is likely to occur to determine when she is most fertile or if she has missed her period and is pregnant. We create a tool to help women track their menstrual cycles to give them a sense of their average cycle length.

Proposed tool. In its simplest form, the tool is a long strip of paper segmented to represent different days. Every seven segments, the strip is highlighted with a different color to indicate the beginning of a week.

Each day the individual does not have their period, she




Figure 4: Average tool as described in Section 3.1.3.

will unroll the strips of paper. On the day the woman does get her period, she will rip the strip of paper at that point to get a segment of dates that represent that one cycle. Once there are paper sections for each of the cycles, these sections are arranged next to each other. This could give a good “eyeball” estimate of the average, and the woman can have a sense of where the average would lie, as seen in the dotted region in Figure 5. The next time, as the woman is unrolling her strip of paper, she will know that as her paper lines up similarly to the previous cycles, she should anticipate to be her period. This could be enhanced with the use of a tool that lines up with the shortest strip and the longest strip in the range of data points and, with springs or rubber bands attached to the middle mark, draws a line between the two to come up with an alternate means of estimating the average. This is seen with the bars and coils in Figure 5.

A different idea for determining averages is the wrapping method. The woman would still continue to unroll the strips each day that she does not get her period. However, she will not rip the strips of paper on the days that she does get her period. Instead, she will have a length of the strip that represents all of the cycles she has had from the beginning. In this example, the woman would have a strip with the length of 167 for her 6 cycles. Then, she will take this long strip and wrap it around two rods, potentially pens or pencils, the number of times she has her cycles. In this example, she has had 6 cycles, so she will wrap the strip tightly six times around the two rods while tracking the end of the 167 day length. She can then hold onto the rods and separate them in a parallel manner in opposite directions until one rod is lined up with the square for the first day and the other rod is lined up with the square for the last day. This should give an even six lengths for the 167 days, which represent the average of the days.

Lessons learned. The strips of paper with date increments can be used for any type of time or date tracking. And the methods of comparing these date strips can be used for any type of averages tracking.

3.2 In the Field Design Iterations

We conducted contextual inquiries in the domains of microfinance and public health in the Greater Accra area in Ghana, including Tema, Dodowa, and Awutu, over a period of three weeks in July 2014. During this time, we observed

#	Date	Deposit	Balance	Receipt #	Officer's Sign.
1					
2					
3					
4					

27	28	29	30	31	1	2	3
					3	3	3
3	4	5	6	7	8	9	6
6		9		12	15		9
10	11	12	13	14	15	16	12
	18	21		24		27	15
							18

#	Date	Deposit	Balance	Overall	Receipt #	Officer's Sign.
1				3		
2				3		
3				6		
4				9		

Figure 5: Susu tracker described in Section 3.1.3.

the organizational settings, user practices, and existing paper artifacts.

3.2.1 Susu Tracker

Use cases. We spent time with a microfinance institution (MFI) in Tema, observing their daily “susu” (savings) operations. Clients deposit a standard amount for 31 days where 1 day is kept by the MFI as commission. These transactions are logged in a standard paper passbook along with the date on which a deposit was made, the amount deposited, and the running total for each cycle. However, the amount taken for commission and the cumulative sum across different cycles is not provided by the passbook. Therefore, when the customer goes for a withdrawal, they are unaware of how much money they have saved up in total and how much of the money they put in was separated for commission. We designed a tool that could give the customer this missing information, specifically geared toward those with low-literacy and arithmetic skills.

Proposed tool. We tested the usability of five design iterations with existing customers. The first iteration is a calendar with a stack of stickers that has the running total, with commission accounted for in the first deposit of a cycle by a red sticker. The final iteration utilizes no stickers, but, in a simple passbook format, has the running total over all of the different cycles already printed out in the overall balance column.

Lessons learned. With the feedback from the users, it was found that the layout that capitalizes on the familiarity of the existing passbook format and required the least amount of additional materials (ie, using only a pencil vs. a separate stack of stickers) is the most effective.

3.2.2 Graph Reader

Use cases. We observed medical practices and conducted interviews with different doctors and nurses at a government district and children’s hospital in the greater Accra region and interviewed medical professors and intermediaries, such as local community health volunteers.

There is a chart that nurses use to determine whether the child is at a healthy weight-age balance based on the WHO regulations of child growth percentiles. This chart also tracks the child’s weight over 24 months, accompanied

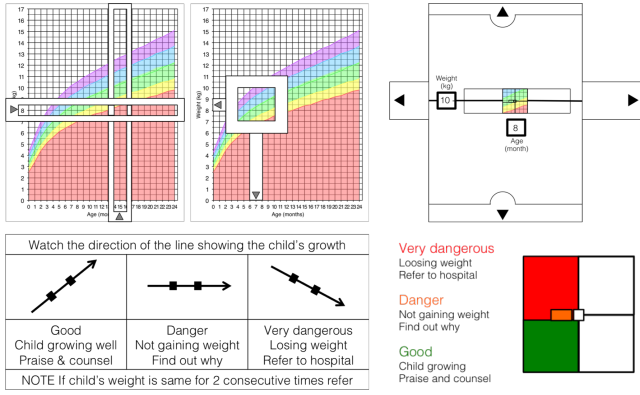


Figure 6: Graph reader in Section 3.2.2.

by a legend to indicate the possible slopes between each of the points, and their implications, as seen in Figure ???. Though this information is relevant to the mothers, the confusing and esoteric presentation makes it inaccessible for those who do not know how to read graphs, and clunky to use for those who could. Therefore, we decided to develop a tool that would help individuals read this graphical information more easily.

Proposed tool. To improve this tool for low-literate populations, the first challenge is finding a way pinpoint the location of the information of interest. Various ideas for tools were developed, including tabs that stretch to the numbers of interest on the graph and position the basic square viewfinder as seen in Figure ???. Another idea was hollow strips that could be placed on the axes to point to a given region and would overlap to pinpoint the specific region in the center, as seen in the second example. However, based on design principles for low-literacy users and feedback from the needs assessment, we decided we not only needed to highlight the important regions but also to cover the information that is not needed. Therefore, we toyed around with overlapping and sliding pieces of paper and developed a tool that uses horizontal and vertical tabs to slide the region of emphasis in place. An example of this tool can be seen later in the paper in Figure 11b.

Another design question was finding ways to compare two points with each other. This could involve something as basic as the slope comparison that the current child and maternal health records use. Or, this could use the viewfinder's space to compare the slope of two points. After feedback from nurses and doctors, the final iteration of the tool includes a viewfinder that compares points based on the quadrants they fall in, as seen in Figure ??.

Lessons learned. We found that the tools that have a narrower field of view and focus on the important information are most effective. Additionally, we find that having a simple viewfinder can help increase the computational feedback provided by the paper graphs. These graph reading tools can be applied to any type of basic line graph.

3.2.3 Flowchart Booklet

Use cases. Through interviews with a reputable non-profit organization engaged in public health interventions and interactions with community health volunteers, we learned that one of the tools used by nurses is a decision tree for diagnosing common ailments. This is derived from the Inte-

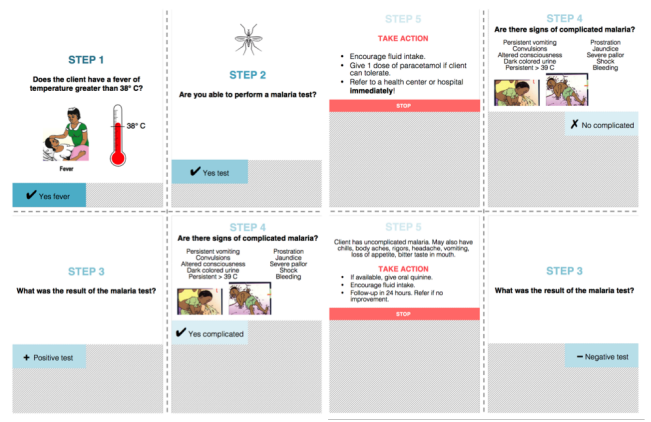


Figure 7: Flowchart booklet in Section 3.2.3.

grated Management of Childhood Illnesses (IMCI), which is a flowchart that instructs the intermediary which questions to ask the caretakers of the patient or which symptoms to check for and how to proceed depending on the results of these inquiries. However, the current representation of the logic for this process is too overwhelming with a lot of detailed information on a single diagram or oversimplifications to accommodate information in a way that is easily represented by a single diagram. Therefore, we decided to create a more efficient tool for presenting this information.

Proposed tool. We started by simplifying the flowchart so that each potential option would be represented by a different section of the booklet. However, we soon found that this would decrease the amount of space that we would have on the page and detract from our ability to properly show all of the information. So we decided to use tabs to demonstrate the options that a user could take at a given point. After demonstrating our first iterations to an audience of doctors and nurses, it was found that the progress from one step to another was not self-explanatory. Even though the users would follow the steps on the booklet, it was difficult to know when to stop. We tried adding arrows to the pages to indicate the flow from one page to the next, but the arrows were unclear when it came to needing to choose which option to flip. Finally, we found that we would make the options more clear but also to append a long red bar that said "STOP" at the end of a given action sequence, as seen in Figure 13a.

Lessons learned. We find that since any "smart" paper tool in a low-resource setting will most likely not come with many user instructions or training sessions, the progression from one step to another has to be fairly clear to the user.

4. DESIGN

Given the immense design space of paper-based tools we necessarily had to limit our scope for this project. We reflected upon the paper tools that we encountered and explored in order to abstract our ideas into basic primitives that could be shared across tools. We then considered the various constraints of paper itself, our target users, and our specific contexts to scope the eventual design of Papyr.

4.1 Paper's Limitations

While the goal of this project was to manipulate paper's

many properties in a way that could augment its computational and visual outputs, we were constantly aware of its limitations:

- Paper has **limited computational feedback** when compared to a computer. As an example, while generating the running total in the susu tracker, we were only able to do so because the intervals were of a fixed value. Therefore, providing a dynamic running total for successive numbers with a variable interval between them would be challenging.
- Paper is also constrained by the **limited amount of physical space** that is available for a specific tool. For instance, the range and the upper limit of numbers in the subtraction band must be directly correlated to the size of the base paper/cardboard sheet the user is willing to work with, and that is not unwieldy and difficult to manage.
- Paper has **limitations on physical manipulation**. Paper sheets, of course, can be folded or placed next to each other, in order to provide some dynamic resizing capabilities, but physical manipulation such as folding also has its limits with respect to the number of folds possible with paper. Paper can provide sliding or pulling/pushing movements (e.g. in the subtraction band as well as the addition/matching tool), but such movements can be awkward or suffer from wear and tear after repeated use. Comparable, low-cost material substitutes, such as ribbons, may be improve certain problems, but we limit our investigation here to paper.

4.2 User Constraints

There were constraints on the tool themselves in terms of how self-explanatory or intuitive they were to the end-user, especially to the low-literate populations that we were targeting. User capabilities can be improved through training and other in-person support, but we limited our investigations to paper tools that required minimal training and intermediation.

During our time in the field and through our interactions with the different users, we also discovered that their goals did not always match the broad goals that we had envisioned for the paper tools. This is not necessarily a user constraint per se, but it did certainly place some checks on the tools we designed and tested in Ghana. For instance, we assumed users wanted more complete and precise results than the susu tracker was providing them. However, in many cases, users were actually satisfied with the current pass-books where they were able to approximate their running total, albeit with some amount of effort. This points to a general inertia towards the learning and implementation of new tools, paper or otherwise. Still, retaining the original look-and-feel of an existing paper tool, something that is more reasonable to accomplish with a new paper tool, like we did with the final iteration of the susu tracker, can help alleviate some of this inertia.

Some of our more sophisticated tools involved too much effort to use, which did not suitably justify the benefit that users were able to glean from them. For instance, the effort invested in using the subtraction band tool may be disproportionate to the benefit that users receive from the tool, especially when their goal is not precision, but is instead a suitably accurate approximation. For similar reasons, while we discussed many paper tool options that would make use of

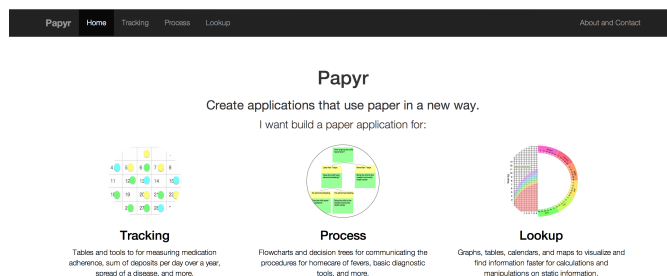


Figure 8: Screenshot of Papyr interface.

such actions as ripping paper to show approximations, layering paper to demonstrate an increase in magnitude through thickness, or even origami to solve more complex problems, we did not end up designing or iterating on many of these ideas. We limited our Papyr prototype to support only the subset of primitives that we felt confident could be of practical use for health and microfinance problems we encountered.

5. PAPYR

Papyr is an interactive system that allows intermediaries working at low-resource organizations to rapidly and automatically generate customized paper tools that give computational and visual feedback. Through a series of questions, the intermediary, as the user of Papyr, specifies a task or tasks. The user is then presented with a set of potential tools generated by Papyr and selects the tool that best suits the resources available and the capabilities of the end user.

In developing Papyr, our main goal is to link the paper primitives that we extracted into an actual usable system. Considering all of the potential types of tasks, the primitives paper could take on, and tools the system could output, we define three different categories of tasks: tracking, process streamlining, and information lookup. The main distinction is whether the information presented is dynamic. Tracking tasks are those that require user input at given time intervals and the resulting information, such as the sum or regularity of a process, is dynamic. Information lookup tasks are ones where the information is static and minimal user interaction is necessary to discover the information. And finally, process tasks are those that require user input to communicate the current, dynamic situation and determine which information would be shown, but the information itself is static.

5.1 Tracking

Tracking encompasses tasks that require periodic user input. This includes tracking the days a patient takes his or her pills over a month to visualize the regularity of medication consumption. Alternatively, it also includes tracking the days the microfinance savings customer makes the required two cedi deposits to calculate cumulative sum over half a year. Other use cases include caloric intake each day over a month, amount of sleep each day, length of menstrual cycle each month, mood each day, journaling of events that occur each day, or spread of disease throughout a community over a month.

These tasks all have an element of time and the information achieved relies on this element of time. In the most

basic form, each of these tasks can be done using a table. The main dimension of this table reflects the change in time, with each row as a single time interval. This varies for each task based on the frequency and the duration for which the task is tracked. Another dimension is the values in each one of the cells. Some tasks require only a simple yes or no of whether the event occurred to determine regularity, such as medication consumption, and other tasks require a numeric value that could contribute to a cumulative sum, such as daily deposits.

In the cases where there is more than one piece of data being tracked over time, each task can be defined separately and then, if the time intervals are compatible, automatically merged by our system into a single table with multiple columns. For example, if the task is to track mood every day, and there are five different possible moods, each of the moods is considered a separate tracking task and merged into a single table with one column for each mood. Users can click a checkbox to disable automatic merging of tasks if this is not desired.

The next challenge is allowing users to input task specifications and parameters in an understandable way and to translate this input to paper tool outputs. We originally proposed to allow the user to string together a series of “per” and “over” phrases for each relevant time and location to specify their task. For example, an input could read “I am tracking the regularity of my medication consumption per day over three months.” However, this method proves to be fairly ambiguous and difficult for both the user and the computer with more complicated inputs. Eventually, the minimal series of questions was developed to determine the dimensions of the tracking task and differentiate between each of the potential outputs.

The resulting system asks the user a series of questions about the tracking task they are proposing, including the values of the input and the frequency and duration of the tracking. The system is interactive—for each question, Papyr iterates on the user’s answers and outputs suggestions, allowing the user to visualize the impact of each additional constraint. For example, after answering the first question, “What are you tracking?”, the user is presented a blank one-column table where only the first row is filled with the task name. The next question asks about the possible values of the tracking task. If the user indicates that they are tracking only whether the event occurs or not, another suggested output is the same one-column table populated with “YES/NO” in each row so that the end-user can circle the corresponding word based on the occurrence of the event. The user can very clearly see the incremental difference between the blank table and the new “YES/NO” table when choosing between the two.

Each of the output suggestions includes the base table that will be used, as well as a set of possible instructions that vary based on the level of literacy of the end user and the cost of any external materials. In the previous example, for the blank table, the user receives potential instructions of writing either “Y” or “N” on each of the lines depending on whether the event occurs, shading in the rows for the time intervals the event occurs, or placing stickers in those rows.

The user then has the opportunity to compare all of the possible tools along with their associated instructions and select the one that matches their specific context. A screen-

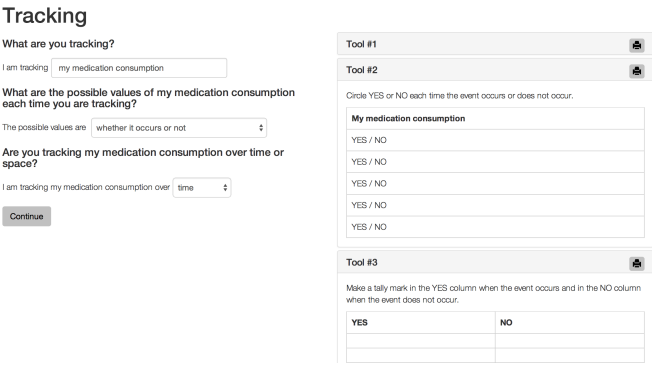


Figure 9: Screenshot of tracking interface.

shot of the Papyr Tracking workflow is seen in Figure 9. The potential table-based outputs are similar to Figure ?? and Figure ?? for “sticker with calendar” and “no sticker with merged tables” outputs respectively.

5.2 Information Lookup

While tracking deals with tasks that require user manipulation and input, other tasks deal with static information. These include tasks such as calculating an individual’s body mass index (BMI) to looking up whether that BMI falls in the percentiles for being over or underweight. Another would be finding out the number of calories for each food item eaten or checking to see if a child’s age and weight are at a healthy balance. The most common forms of static input and output are calculations and information lookups. And in reality, calculations can be construed as information lookups because the answers to a calculation could be enumerated and placed in a lookup table. As with the tracking baseline representation, our baseline representation for lookup is in the form of tables.

The dimensions of these tables are essentially the inputs (i.e. a lookup key) and the outputs (i.e. information the end user is looking up), which are likely numeric or can be reduced to being numerical i.e. categorical). In a basic lookup table, all of the inputs would be in the first columns and the associated outputs would be in the columns that follow, with all possible combinations in each of the rows. This basic layout can be seen in Figure 10a.

However, there are many cases in which there are more appropriate ways to represent this data than this large, cumbersome table. In a more specific example, where there are two inputs and one output, the inputs and outputs could be arranged such that one input is the columns and the other is the rows and the outputs as the corresponding intersecting cells. This can be visualized in Figure 10b. In another case where there are potentially three inputs and one output, a similar row-column structure would work for the presentation of two of the three inputs and the final input would be factored into the separation of the table into tables based on different values for that input. This can be visualized in Figure 10c. And with even larger numbers of inputs, the split could be across different pages, with different colors, and other methods that engage the dimensionality of paper primitives as discussed in the abstraction section.

For low-dimensional lookups, tables are relatively easy to use. For higher-dimensional lookup tasks, alternative visu-

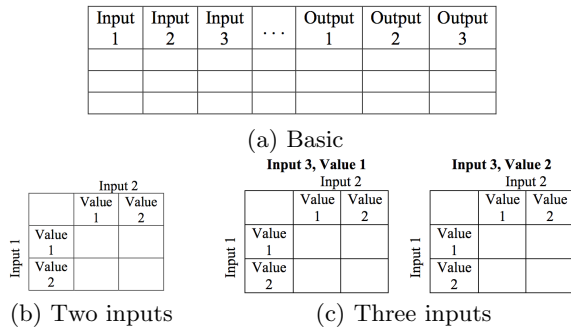


Figure 10: Different arrangements of lookup tables.

alizations are beneficial. A clock or calendar could be used to visualize information that has to do with time or date. Additionally, information could also be presented on a map. Due to lack of space, the two specialized lookup outputs we present here are graph representations and calendar year calculations.

Graph representations cover data that has two numerical dimensions as inputs and a numerical or categorical output. For example, if someone knows their height and weight as numerical inputs, they can figure out whether they fall in the category of under, normal, or overweight based on a graph with those categories as differently colored regions.

Papry suggests graph enhancements to help end users with understanding the content of the chart or encourage interactivity. These improvements are drawn from the ideas and feedback collected from the graph reading tool we proposed. The first improvement includes tabs and cutouts to help narrow the end user's focus on the relevant region of the graph. For graphs with two axes, this is a vertical sliding cover to isolate based on the y-axis and a horizontal sliding tab to isolate the x-axis. The user can elect to use these components together or separately. This can be seen in Figure 11 which shows the output of a graph tool that has the sliding tabs that pinpoint the relevant information to help block out the rest of the graph.

A second enhancement is a "smart viewfinder" that guides the end user in doing some basic approximations. For example, sectioning the "smart viewfinder" into different quadrants in relation to the midpoint helps to determine if the slope of the line from the previous point is positive or negative (e.g. Figure ??). A notched "smart viewfinder" could help the end-user compare the distance between the two points, whether it is comparing the sizes of bars on a bar graph or the amount needed to increase or decrease to be in a different category or region.

In order to build these graph-based information lookup tools, users will be asked to input the data that they want to work with, and a graph will be built. Then, users will be presented with the pros and cons of each graph improvement and asked which enhancements they want, whether it is the vertical or horizontal tabs or the viewfinder.

Calendar year calculations are a specific type of calculation for which we developed a paper tool. This tool is made of two concentric circles, the smaller layered on top of the larger. Users can specify the relative number of days between one event and another, provided the difference is within the 365 day range of a year. Then, by spinning the inner circle to line up with a date on the outer ring, the end

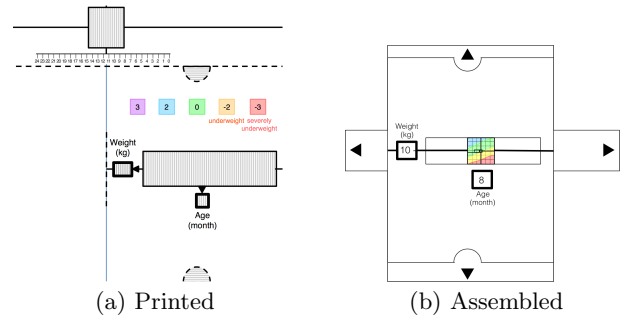


Figure 11: Output of the graph generation tool printed and assembled.




Figure 12: Example of calendar year calculation tool.

user can calculate the date of the next event. For example, gestation occurs approximately 280 days after the last menses. A tool for this aligns the 0 day mark with the date of the last menses, the 14 day mark with conception, and the 280 day mark lines up with the predicted gestation day. This sample tool is seen in Figure 12. To generate this tool for other calculations, the Papry system takes as an input an indication of which days are matched with which important event in the year range.

5.3 Process Streamlining

A specific case of information lookup is process driven—though it has static information, the relevance of this information depends on the specifics of the situation. This information could be presented using a table, with each of the potential factors that contributes to the classification of a given situation as an input and the resulting instructions as an output. However, this would result in a fairly sparse table or many sets of sparse tables separated on different pages or sections. More commonly, this type of information is presented in a flowchart or decision tree. The tool produced by the process module of Papry can be used to depict this logic, including, but not limited to, questions community health volunteers could use to diagnose a patient, decision diagrams to help a customer choose a savings scheme, or checklists for procedures that change based on the outcome of a step.

In the Papry system, the user can specify the logic of the process in flowchart form. The flowchart is a directed, acyclic graph whose nodes are the "steps" or instructions the end user should follow and whose edges are "options" or

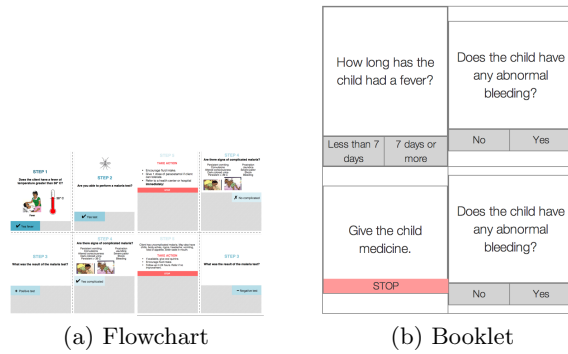


Figure 13: Process logic depicted as a flowchart and as a booklet.

logical connections from one step to another.

For example, if the user is looking to create a tool to help parents with the home treatment of fever in their children, the user would first create a step with the question of “How long has the child had a fever?” From this question, there would be two possible paths that can be taken. The user would create another step connected to the first step with the option “less than seven days” and the instruction to “bring the child to the nearest community health center.” For the other possible answer to the question, the user would create another step that is connected to the first question with the option “more than seven days” and with the question “Does the child have any abnormal bleeding?”. The user would then continue this process in specifying the answers to this question and continue until no questions remain and all paths end in instruction. This example is presented in Figure 13a.

Then, the system converts this logic into a booklet. Each page of the booklet will present one question from the flowchart and the potential answer options associated with it as tabs below. The end-user can flip the tab over of the option that they select and continue with either the question or the instruction on the page. An example of the output of the Papyr system is shown in Figure 13b.

Some of the challenges for this workflow in Papyr include figuring out the order of the steps and instructions. In order to incorporate the feedback from the needs assessment that there should be no blank pages, balanced with the need to show visible instructions, whether it is all possible options at a given time or a “STOP” instruction, our resulting algorithm for adapting the process logic from a flowchart to a booklet ended up being a depth first traversal that highlighted all nodes, or steps, that were accessed, even those that had been visited before. At each node, the algorithm would access the children of the node and append these options as tabs at the bottom of the page but if there were no children, append a “STOP”. Another challenge was determining the best way to fit all of the instructions on a page, as per the “screen real estate” constraint of paper. The default presentation of information is using a fourth of an A4 sheet of paper. However, since there seemed to be varying preferences from the feedback, the user also has the option to adjust the size of each of the pages as well and our system uses this to parameterize the output.

5.4 Implementation

The system is implemented as a web-based system with HTML/CSS and Javascript built utilizing the jQuery, Twitter Bootstrap, and D3 frameworks in 2134 lines of code.

6. DISCUSSION

The motivation driving this study is the fact that paper is abundant, affordable, deeply embedded in existing workflow practices and personal spaces, and familiar, especially when compared to digital devices that are yet to achieve the same level of adoption and use in the developing world. Furthermore, as we demonstrate in this work, paper is capable of providing augmented computational and visual feedback to achieve some of the everyday tasks that computers currently do. Smart paper can perform basic arithmetic tasks, provide regularity and frequency information through visual feedback, and streamline information dissemination.

In general, when aiming to solve a problem in a platform-agnostic way (should I use paper or technology?), the design motivations, principles, and applications remain broadly consistent (i.e. without considering technology, all the other design considerations remain the same). Therefore, while working within a given scope, that is the degree of computational or visual feedback required to solve a problem, we encourage a more rigorous and dynamic use of paper to resolve some of these problems, especially in low-resource settings. We believe that the fervent drive to use ICTs to solve a broad range of problems could be scaled back and reconsidered. To avoid the risk of romanticizing paper, we recognize that paper’s computational power cannot match that of a computer’s. However, paper can certainly be improved over its existing uses as we show specifically in the context of low-resource organizations in public health and microfinance.

Papyr is a supporting infrastructure that automates the creation of the paper tools. Papyr is unlike existing paper computing applications that almost always reduce paper’s presence in these applications, diluting some of its more valuable affordances, or replacing it altogether by mimicking its properties in a surrogate (and “sexy”) technology artifact. Instead, Papyr leverages technology in the design and creation processes, but the output completely preserves the affordances of paper for the end user.

7. CONCLUSIONS AND FUTURE WORK

This paper describes an exploration into the design of “smart” paper tools in the context of developing regions. We described the process of brainstorming, designing, assessing, and developing a paper-technology infrastructure that automates the production and distribution of smart paper tools. We explored existing paper tool use in healthcare and microfinance in Ghana to understand how paper is currently being used by low-resource organizations to track and disseminate information. From our needs assessment we iterated on several potential designs of enhanced paper-based tools to push the boundaries of paper. We designed a smart paper generation tool that allows intermediaries to specify tasks to quickly and automatically construct potential paper-based solutions that do not require ICTs at the point of use. We describe the design and implementation of our “smart paper” prototype, Papyr and demonstrate how it can be used to produce smart paper outputs that preserve the affordances of paper for the end user.

There are many features that we are still incorporating into Papyr and even more areas of improvement. We believe that there is ample potential to leverage paper's tangibility in future work. We plan to explore many of the ideas we mentioned here including: ripping, stacking, layering, or complex folding (origami) with paper. There is also potential to improve the specific modules in Papyr. For instance, future iterations of the tracking workflow will include improved abstractions and optimizations from the given tables, so as to better handle more inputs across different tasks, including optimizations of using maps to track the events across geographical space. In the same vein, future versions of the information lookup workflow will produce more specialized tools including editable calendars and maps. Finally, a future optimization on process streamlining will allow the inclusion of images as components of the steps, for both the questions as well as the instructions and options to better accommodate low-literate individuals. We plan to provide Papyr as a free open-source platform for immediate use by low-resource organizations.

8. ACKNOWLEDGEMENTS

We thank the participants of our studies for their patience, helpfulness, and valuable feedback.

9. REFERENCES

- [1] Y. Anokwa, C. Hartung, W. Brunette, G. Borriello, and A. Lerer. Open source data collection in the developing world. *Computer*, 42(10):97–99, 2009.
- [2] Anoto. Digital pen and paper technology. *Anoto AB*.
- [3] C. D. Chandler, G. Lo, and A. K. Sinha. Multimodal theater: Extending low fidelity paper prototyping to multimodal applications. In *CHI'02 Extended Abstracts on Human Factors in Computing Systems*, pages 874–875. ACM, 2002.
- [4] K. Chen, A. Kannan, Y. Yano, J. Hellerstein, and T. Parikh. March 2012.
- [5] M. Coelho, L. Hall, J. Berzowska, and P. Maes. Pulp-based computing: a framework for building computers out of paper. In *CHI'09 Extended Abstracts on Human Factors in Computing Systems*, pages 3527–3528. ACM, April 2009.
- [6] N. Dell, N. Breit, J. O. Wobbrock, and G. Borriello. Improving form-based data entry with image snippets. In *Proceedings of the 2013 Graphics Interface Conference*, pages 157–164. Canadian Information Processing Society, 2013.
- [7] N. Dell, J. Crawford, N. Breit, T. Chaluco, A. Coelho, J. McCord, and G. Borriello. Integrating odk scan into the community health worker supply chain in mozambique. In *Proceedings of the Sixth International Conference on Information and Communication Technologies and Development: Full Papers-Volume 1*, pages 228–237. ACM, 2013.
- [8] M. Fahdhy and V. Chongsuvivatwong. Evaluation of world health organization partograph implementation by midwives for maternity home birth in medan, indonesia. *Midwifery*, 21(4):301–310, 2005.
- [9] A. Fawole, K. Hunyinbo, and D. Adekanle. Knowledge and utilization of the partograph among obstetric care givers in south west nigeria. *African journal of reproductive health*, 12(1):22–29, 2009.
- [10] A. Gomes, A. Nesbitt, and R. Vertegaal. Morephone: A study of actuated shape deformations for flexible thin-film smartphone notifications. In *Proceedings of ACM CHI'13 Conference on Human Factors in Computing Systems*, pages 583–592. ACM Press, 2013.
- [11] A. Gomes and R. Vertegaal. Paperfold: A shape changing mobile device with multiple reconfigurable electrophoretic magnetic display tiles. In *Extended Abstracts of ACM CHI'14 Conference on Human Factors in Computing*. ACM Press, 2014.
- [12] W. Johnson, H. Jellinek, L. K. Jr, R. Rao, and S. Card. Bridging the paper and electronic worlds: the paper user interface. In *Proceedings of the INTERACT'93 and CHI'93 conference on Human factors in computing systems*, pages 507–512. ACM, 1993.
- [13] F. Kaplan and P. Jermann. Papercomp 2010: first international workshop on paper computing. In *Proceedings of the 12th ACM international conference adjunct papers on Ubiquitous computing-Adjunct*, pages 507–510. ACM, 2010.
- [14] W. H. O. Maternal, C. Health, and F. Planning. The partograph: a managerial tool for the prevention of prolonged labour (vol. 1). 1988.
- [15] W. Newman and P. Wellner. A desk supporting computer-based interaction with paper documents. In *Proceedings of the SIGCHI conference on Human factors in computing systems*. ACM, 1992.
- [16] T. Parikh, S. Javid, K. Ghosh, and K. Toyama. April 2006.
- [17] K. Pettersson, M. Svensson, and K. Christensson. Evaluation of an adapted model of the world health organization partograph used by angolan midwives in a peripheral delivery unit. *Midwifery*, 16(2):82–88, 2000.
- [18] T. Pietrzak, S. Malacria, and Ã. Lecolinet. What do u-note? an augmented note taking system for the classroom. In *Papercomp Workshop*, 2010.
- [19] A. Ratan, K. Toyama, S. Chakraborty, S. Ooi, K. Koenig, M. Chitnis, and M. Phiong. December 2010.
- [20] A. Sellen and R. Harper. Paper as an analytic resource for the design of new technologies. In *Proceedings of the ACM SIGCHI Conference on Human factors in computing systems*, pages 319–326. ACM, 1997.
- [21] A. Sellen and R. Harper. *The myth of the paperless office*. MIT Press, 2002.
- [22] B. Signer and M. C. Norrie. Paperpoint: a paper-based presentation and interactive paper prototyping tool. In *Proceedings of the 1st international conference on Tangible and embedded interaction*, pages 57–64. ACM, 2007.
- [23] C. Snyder. *Paper prototyping: The fast and easy way to design and refine user interfaces*. Newnes, 2003.
- [24] A. Tarun, P. Wang, A. Girouard, P. Strohmeier, D. Reilly, and R. Vertgaal. Papertab: Tablets as thin and flexible as paper. In *Extended Abstracts of ACM CHI'13 Conference on Human Factors in Computing Systems*, pages 2881–2882. ACM Press, 2013.
- [25] A. Thomas, D. McKay, and M. Cutlip. A nomograph method for assessing body weight. *The American*

- journal of clinical nutrition*, 29(3):302–304, 1976.
- [26] A. Umezulike, H. Onah, and J. Okaro. Use of the partograph among medical personnel in enugu nigeria. *International Journal of Gynecology and Obstetrics*, 65(2):203–5, 1999.
 - [27] H. Underwood. Partopen: enhancing the partograph with digital pen technology. In *CHI’12 Extended Abstracts on Human Factors in Computing Systems*, pages 1393–1398. ACM, 2012.
 - [28] H. Underwood, S. R. Sterling, and J. K. Bennett. The partopen in training and clinical use-two preliminary studies in kenya. In *HEALTHINF*, pages 112–121, 2013.
 - [29] S. Van Wart, K. J. Tsai, and T. Parikh. Local ground: a paper-based toolkit for documenting local geo-spatial knowledge. In *Proceedings of the First ACM Symposium on Computing for Development*, page 11. ACM, 2010.
 - [30] P. Wellner. Interacting with paper on the digitaldesk. *Communications of the ACM*, pages 87–96, 1993.